

# CloudFormation



- Currently, we have been doing a lot of manual work.
- All this manual work will be very tough to reproduce:
  - In another region
  - In another AWS account
  - Within the same region if everything was deleted.
- Would it be a great if we had a code that would deploy everything in a code?
- That code would be deploy and create / update / delete our infrastructure.

## What is CloudFormation?

- CloudFormation is a **declarative way** of outlining your AWS infrastructure for any resources (most of them are supported).
- For example, within a CloudFormation template, you say:
  - I want a security group
  - I want two EC2 instances using the Security Groups
  - I want two EIP s for the EC2 machines
  - I want an S3 bucket
  - I want a load Balancer (ELB) in front of these machines.
- Then CloudFormation creates those for you, in the **right order**, with the **exact configuration** that you specify.

## Benefits of CloudFormation

- **Infrastructure as a Code:**
  - No resource are manually created, which is excellent for control.
  - The code can be version controlled for example using git.
  - Changes to the infrastructure are reviewed through code.
- **Cost:**
  - Each resource within the stack is staged with an identifier so you can easily see how much a stack costs you.
  - You can estimate the cost of your resources using CloudFormation template.
  - Savings strategy: In Dev, you can make an automation to delete the template at 5 pm and recreate it at 8 AM safely.
- **Productivity:**
  - Ability to destroy and re-create an infrastructure on the cloud on the fly.
  - Automated generation of Diagram for your templates!
  - Declarative programming ( no need to figure out ordering and orchestration)

- **Separation of concern: create many stacks for many apps, and many layers. Example:**
  - VPC stack
  - Network Stacks
  - App stacks
- **Do not re-invent the wheels:**
  - Leverage existing templates on the web!
  - Leverage the documentation.

## How CloudFormation Works

- Templates has to be uploaded in S3 and then referenced in the CloudFormation.
- To update a template, we can not edit previous ones. We have to re-upload a new version of the template to AWS
- Stacks are identified by a name
- Deleting a stack deletes every single artifact that was created by CloudFormation.

## Deploying the CloudFormation templates:

### 1. Manual way:

- Editing templates in the CloudFormation Designer.
- Using the console to input parameters, etc

### 2. Automated way:

- Editing templates in a YAML file
- Using the AWS CLI (Command Line Interface) to deploy the templates.
- Recommended way when you fully want to automate your flow

## CloudFormation Building Blocks

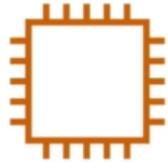
### Template components

1. **RESOURCES:** your AWS resources declared in the template (Mandatory)
2. **PARAMETERS:** the dynamic inputs for your template
3. **MAPPINGS:** the static variables for your templates
4. **OUTPUTS:** Reference to what has been created
5. **CONDITIONALS:** List fo conditions to preform resource creation
6. **METADATA:**

### Template helpers:

1. **Reference**
2. **Functions**

# Introduction to CF.



We are going to create a simple EC2 instance.  
Then we are going to create to add an EIP to it.  
We are going to add two Security Groups to it.

EC2 Instance

- 1) Go to the CloudFormation and press Create Stack. This will be our first stack.
- 2) Our template is ready in the file we will just go and upload that file. Next

Step 1  
**Specify template**

Step 2  
**Specify stack details**

Step 3  
**Configure stack options**

Step 4  
**Review**

## Create stack

### Prerequisite - Prepare template

#### Prepare template

Every stack is based on a template. A template is a JSON or YAML file that contains configuration information about the AWS resources you want to include in the stack.

Template is ready

Use a sample template

Create template in Designer

### Specify template

A template is a JSON or YAML file that describes your stack's resources and properties.

#### Template source

Selecting a template generates an Amazon S3 URL where it will be stored.

Amazon S3 URL

Upload a template file

#### Upload a template file

Choose file

0-just-ec2.yaml

JSON or YAML formatted file

[View in Designer](#)

- 3) Give a Stack name.
- 4) You can give a stack name and press create for now.

### Specify stack details

#### Stack name

Stack name

firs

Stack name can include letters (A-Z and a-z), numbers (0-9), and dashes (-).

#### Parameters

Parameters are defined in your template and allow you to input custom values when you create or update a stack.

#### No parameters

There are no parameters defined in your template

[Cancel](#)

[Previous](#)

[Next](#)

Stacks (1)		
Filter by stack name		
Active <input checked="" type="checkbox"/> View nested < 1 >		
first-stack	2020-05-06 18:56:12 UTC+0100	
		CREATE_COMPLETE

Events (5)				
Search events				
Timestamp	Logical ID	Status	Status reason	
2020-05-06 18:56:50 UTC+0100	first-stack		CREATE_COMPLETE	-
2020-05-06 18:56:49 UTC+0100	MyInstance		CREATE_COMPLETE	-
2020-05-06 18:56:17 UTC+0100	MyInstance		CREATE_IN_PROGRESS	Resource creation Initiated
2020-05-06 18:56:15 UTC+0100	MyInstance		CREATE_IN_PROGRESS	-
2020-05-06 18:56:12 UTC+0100	first-stack		CREATE_IN_PROGRESS	User Initiated

Stacks (1)		
Filter by stack name		
Active <input checked="" type="checkbox"/> View nested < 1 >		
first-stack	2020-05-06 18:56:12 UTC+0100	
		CREATE_COMPLETE

## first-stack

[Delete](#) [Update](#) [Stack actions ▾](#) [Create stack ▾](#)

- [Stack info](#)
- [Events](#)
- [\*\*Resources\*\*](#)
- [Outputs](#)
- [Parameters](#)
- [Template](#)
- [Change sets](#)

### Resources (1)

Logical ID	Physical ID	Type	Status	Status reason
MyInstance	<a href="#">i-005d32e2ee0eb698</a>	AWS::EC2::Instance		CREATE_COMPLETE

The “Resource” section will show what kind of resources has been created so far.

Stacks (1)		
Filter by stack name		
Active <input checked="" type="checkbox"/> View nested < 1 >		
first-stack	2020-05-06 18:56:12 UTC+0100	
		CREATE_COMPLETE

## first-stack

[Delete](#) [Update](#) [Stack actions ▾](#) [Create stack ▾](#)

- [Stack info](#)
- [Events](#)
- [Resources](#)
- [Outputs](#)
- [Parameters](#)
- [\*\*Template\*\*](#)
- [Change sets](#)

### Template

[View in Designer](#)

```

---
Resources:
  MyInstance:
    Type: AWS::EC2::Instance
    Properties:
      AvailabilityZone: us-east-1a
      ImageId: ami-a4c7edb2
      InstanceType: t2.micro
  
```

Will show us simple “Template” file for what we have created so far.

So far this is a vary simple setup and we have created the simple EC2 instance.

In the “**Events**” section - You can see all the progress what you have been done so far.

# CloudFormation Update and Delete Stack

! 0-just-ec2.yaml ×

```
cloudformation > ! 0-just-ec2.yaml > {} Resources > {} MyInstance > {} Properties > ImageId
```

```
1 ---  
2 Resources:  
3   MyInstance:  
4     Type: AWS::EC2::Instance  
5     Properties:  
6       AvailabilityZone: us-east-1a  
7       ImageId: ami-a4c7edb2  
8       InstanceType: t2.micro  
9
```

On the second template what we will do is. We will add

- EIP
- Two SG
- And EC2

! 1-ec2-with-sg-eip.yaml ×

```
! 1-ec2-with-sg-eip.yaml > {} Resources > {} ServerSecurityGroup > {} Properties > [ ] SecurityGroupIngress > {} 0 > # ToPo  
1 ---  
2 Parameters:  
3   SecurityGroupDescription:  
4     Description: Security Group Description  
5     Type: StringType  
6  
7 Resources:  
8   MyInstance:  
9     Type: AWS::EC2::Instance  
10    Properties:  
11      AvailabilityZone: us-east-1a  
12      ImageId: ami-a4c7edb2  
13      InstanceType: t2.micro  
14      SecurityGroups:  
15        - !Ref SSHSecurityGroup  
16        - !Ref ServerSecurityGroup  
17  
18    # an elastic IP for our instance  
19    MyEIP:  
20      Type: AWS::EC2::EIP  
21      Properties:  
22        InstanceId: !Ref MyInstance  
23  
24    # our EC2 security group  
25    SSHSecurityGroup:  
26      Type: AWS::EC2::SecurityGroup  
27      Properties:  
28        GroupDescription: Enable SSH access via port 22  
29        SecurityGroupIngress:  
30          - CidrIp: 0.0.0.0/0  
31          | FromPort: 22  
32          | IpProtocol: tcp  
33          | ToPort: 22  
34  
35    # our second EC2 security group  
36    ServerSecurityGroup:  
37      Type: AWS::EC2::SecurityGroup  
38      Properties:  
39        GroupDescription: !Ref SecurityGroupDescription  
40        SecurityGroupIngress:  
41          - IpProtocol: tcp  
42          | FromPort: 80  
43          | ToPort: 80  
44          | CidrIp: 0.0.0.0/0  
45          - IpProtocol: tcp  
46          | FromPort: 22  
47          | ToPort: 22  
48          | CidrIp: 102.168.1.1/22
```

So in the above we have created an EC2 instance simple template in YAML file. But what we are going to do is we will **update the stack** with the different template.

AWS	Services	Resource Groups		stephane @ datacumulus-cour...	N. Virginia	Support
CloudFormation > Stacks						
<b>Stacks (1)</b>						
<input type="text"/> Filter by stack name	<input type="button"/>	<input type="button"/> Delete	<input type="button"/> Update	<input type="button"/> Stack actions	<input type="button"/> Create stack	
<input checked="" type="checkbox"/> Active	<input type="checkbox"/> View nested	< 1 >	<input type="button"/>			

**1) To Update a Stack - you need to go back and select the stack and press “update”**

CloudFormation > Stacks > first-stack > Update stack

Step 1  
Specify template

Step 2  
Specify stack details

Step 3  
Configure stack options

Step 4  
Review

**Update stack**

**Prerequisite - Prepare template**

Prepare template  
Every stack is based on a template. A template is a JSON or YAML file that contains configuration information about the AWS resources you want to include in the stack.

Use current template    Replace current template    Edit template in designer

**Cancel** **Next**

**2) There will be a 3 options and we will choose to “Replace current template”**

**Specify template**

A template is a JSON or YAML file that describes your stack's resources and properties.

**Template source**  
Selecting a template generates an Amazon S3 URL where it will be stored.

Amazon S3 URL    Upload a template file

**Upload a template file**

**Choose file**  1-ec2-with-sg-eip.yaml

JSON or YAML formatted file

S3 URL: <https://s3-external-1.amazonaws.com/cf-templates-137vmwb32gjrt-us-east-1/2020127Ub4-1-ec2-with-sg-eip.yaml>  View in Designer

**Cancel** **Next**

We will upload a new file and as you can see the new file has been uploaded to the S3 but we choose our file and press Next.

## Change set preview

### Changes (4)

Search changes

< 1 >

Action	Logical ID	Physical ID	Resource type	Replacement
Add	MyEIP	-	AWS::EC2::EIP	-
Modify	MyInstance	i-005d632e2ee0eb698	AWS::EC2::Instance	True
Add	SSHSecurityGroup	-	AWS::EC2::SecurityGroup	-
Add	ServerSecurityGroup	-	AWS::EC2::SecurityGroup	-

[Cancel](#)

[Previous](#)

[View change set](#)

[Update stack](#)

In the Review section we get the all the changes that will be done to our existing template. For the Update and it will show us what changes will be done to our template.

Press Update Stack. And as we expected our resources has been created. And it will update our previous stack.

### Deletion Process:

You can either delete manually every resource one by one or what you can do is the right way is delete the stack itself. And it will delete the whole resources that has been created within the stack. And CloudFormation knows what to delete in the right order.

CloudFormation > Stacks

### Stacks (1)

[C](#)

[Delete](#)

[Update](#)

[Stack actions ▾](#)

[Create stack ▾](#)

Filter by stack name

Active ▾

View nested

< 1 >

Stack name	Status	Created time	Drift status	Last drift check time
first-stack	UPDATE_COMPLETE	2020-05-06 18:56:12 ...	NOT_CHECKED	-

# YAML Crash Course

```

1 invoice:      34843
2 date   :     2001-01-23
3 bill-to:
4   given  :  Chris
5   family : Dumars
6   address:
7     lines: |
8       458 Walkman Dr.
9       Suite #292
10      city   : Royal Oak
11      state  : MI
12      postal : 48046
13 product:
14   - sku        : BL394D
15     quantity   : 4
16     description: Basketball
17     price      : 450.00
18   - sku        : BL4438H
19     quantity   : 1
20     description: Super Hoop
21     price      : 2392.00

```

CloudFormation uses both YAML and JSON formats but we will go through the YAML in our case because it is little bit easier to use and understand.

**Key value Pairs** - invoice : 34843 line1

**Support Arrays** - line 14 usually it starts with “-” sign sku.

**Nested Objects** - line3 and within it we have nested objects like line4,5,6, and within line 6 we have nested objects line line 7,10 etc

**Multi line Strings** - line 7 “|” is the multi line support.

So far we have created the EC2 instance in the YAML format.

## PARAMETERS

- Parameters are a way to provide inputs to your AWS CloudFormation template.
- Why you need to define parameters:
  - You want to **reuse** your templates across the company
  - Some inputs can not be determined ahead of time
- Parameters are extremely powerful, controlled and can prevent errors from happening in your templates **thanks to types**.

### When you should use the Parameters?

You can use Parameters

if the CloudFormation resource configuration likely to change in the future? If yes then use parameters.

You don't have to re-upload the template to change it is content.

#### Parameters:

##### SecurityGroupDescription:

Description: Security Group Description  
(Simple parameter)

Type: String

# Parameters Settings

Parameters can be controlled by all these settings:

- Type:
  - String
  - Number
  - CommaDelimitedList
  - List<Type>
  - AWS Parameter (to help catch invalid values – match against existing values in the AWS Account)
- Description
- Constraints
  - ConstraintDescription (String)
  - Min/MaxLength
  - Min/MaxValue
  - Defaults
  - AllowedValues (array)
  - AllowedPattern (regexp)
  - NoEcho (Boolean)

## How to Reference a Parameter

- ✓ The **Fn::Ref** function can be leveraged to reference parameters.
- ✓ Parameters can be used anywhere in a template
- ✓ The short stand for this in YAML is **!Ref**
- ✓ The function can also reference other elements within the template.

```
! 0-just-ec2.yaml      ! 1-ec2-with-sg-eip.yaml *
```

```
1 ---  
2 Parameters:  
3   SecurityGroupDescription:  
4     Description: Security Group Description  
5     Type: String  
35 # our second EC2 security group  
36 ServerSecurityGroup:  
37   Type: AWS::EC2::SecurityGroup  
38   Properties:  
39     GroupDescription: !Ref SecurityGroupDescription  
40     SecurityGroupIngress:  
41       - IpProtocol: tcp  
42         FromPort: 80  
43         ToPort: 80  
44         CidrIp: 0.0.0.0/0  
45       - IpProtocol: tcp  
46         FromPort: 22  
47         ToPort: 22  
48         CidrIp: 192.168.1.1/32
```

In the Example below as you can see we, defined our parameters Line2. It is simple description parameter. And in the line39 we Reference it to the Parameter using **!Ref** function. But you can use this **!Ref** function to reference to other resource in the Resource itself. Like EIP resource to the EC2 with **!Ref**

# RESOURCES

- Resources are the core of the CloudFormation templates (Mandatory)
  - They represent the different AWS components that will be created and configured
  - Resources are declared and can reference each other.
  - AWS figures out creation, updates and deletes of resources for us.
  - There are over 224 types of resources (!)
  - Resources types identifiers are of the form:

**AWS::aws-product-name::data-type-name**

To look for the resource you can always google it and get it from there.

```
! 1-ec2-with-sg-eip.yaml X
! 1-ec2-with-sg-eip.yaml > {} Resources > {} ServerSecurityGroup > {} Properties > [ ] SecurityGroupIngress > {} 0 > # ToPo
1 ---  
2 Parameters:  
3   SecurityGroupDescription:  
4     Description: Security Group Description  
5     Type: StringType  
6  
7 Resources:  
8   MyInstance:  
9     Type: AWS::EC2::Instance  
10    Properties:  
11      AvailabilityZone: us-east-1a  
12      ImageId: ami-a4c7edb2  
13      InstanceType: t2.micro  
14      SecurityGroups:  
15        - !Ref SSHSecurityGroup  
16        - !Ref ServerSecurityGroup  
17  
18   # an elastic IP for our instance  
19   MyEIP:  
20     Type: AWS::EC2::EIP  
21     Properties:  
22       InstanceId: !Ref MyInstance  
23  
24   # our EC2 security group  
25   SSHSecurityGroup:  
26     Type: AWS::EC2::SecurityGroup  
27     Properties:  
28       GroupDescription: Enable SSH access via port 22  
29       SecurityGroupIngress:  
30         - CidrIp: 0.0.0.0/0  
31           FromPort: 22  
32           IpProtocol: tcp  
33           ToPort: 22  
34  
35   # our second EC2 security group  
36   ServerSecurityGroup:  
37     Type: AWS::EC2::SecurityGroup  
38     Properties:  
39       GroupDescription: !Ref SecurityGroupDescription  
40       SecurityGroupIngress:  
41         - IpProtocol: tcp  
42           FromPort: 80  
43           ToPort: 80  
44           CidrIp: 0.0.0.0/0  
45         - IpProtocol: tcp  
46           FromPort: 22  
47           ToPort: 22  
48           CidrIp: 192.168.1.1/32
```

For example in the Resource block we give the small indentation and we defined all the resources under the Resource block.

**Type** and **Properties**  
has to be defined.

# MAPPING

- Mapping are fixed variables within your CloudFormation Templates
- They are handy to differentiate between different environments (dev vs prod), regions (AWS regions), AMI types etc.
- All the values are hardcoded within the template

Example:

```
Mappings:  
  Mapping01:  
    Key01:  
      Name: Value01  
    Key02:  
      Name: Value02  
    Key03:  
      Name: Value03
```

```
RegionMap:  
  us-east-1:  
    "32": "ami-6411e20d"  
    "64": "ami-7a11e213"  
  us-west-1:  
    "32": "ami-c9c7978c"  
    "64": "ami-cfc7978a"  
  eu-west-1:  
    "32": "ami-37c2f643"  
    "64": "ami-31c2f645"
```

## When would you use mapping vs parameters?

- Mapping are great when you know in advance all the values that can be taken and that they can be deducted from variables such as:
  - Region
  - Availability Zone
  - AWS Account
  - Environment (dev vs prod)
- They allow safer control over the template.
- Use parameters when the values are really user specific.

## Fn::FindInMap (Accessing Mapping Values)

We use Fn::FindInMap to return a named value from a specific key

```
!FindInMap [MapName, TopLevelKey, SecondLevelKey]
```

## Concept: Pseudo Parameters

- AWS offers us pseudo parameters in any CloudFormation template.
- These can be used at any time and are enabled by default

Reference Value	Example Return Value
AWS::AccountId	1234567890
AWS::NotificationARNs	[arn:aws:sns:us-east-1:123456789012:MyTopic]
AWS::NoValue	Does not return a value.
AWS::Region	us-east-2
AWS::StackId	arn:aws:cloudformation:us-east-1:123456789012:stack/MyStack/1c2fa620-982a-11e3-aff7-50e2416294e0
AWS::StackName	MyStack

```
AWSTemplateFormatVersion: "2010-09-09"  
Mappings:  
  RegionMap:  
    us-east-1:  
      "32": "ami-6411e20d"  
      "64": "ami-7a11e213"  
    us-west-1:  
      "32": "ami-c9c7978c"  
      "64": "ami-cfc7978a"  
    eu-west-1:  
      "32": "ami-37c2f643"  
      "64": "ami-31c2f645"  
Resources:  
  myEC2Instance:  
    Type: "AWS::EC2::Instance"  
    Properties:  
      ImageId: !FindInMap [RegionMap, !Ref "AWS::Region", 32]  
      InstanceType: m1.small
```

# OUTPUTS

- The Outputs section declares Optional outputs values that we can import into other stacks (if you export then first)!
- You can also view the outputs in the AWS Console or in using AWS CLI
- They are very useful example if you define a network CloudFormation and outputs the variables such as VPC ID and SubnetIDs
- It is the best way to perform some collaboration cross stack, as you let expert handle their own part of the stack
- You can not delete a CloudFormation Stack if it is outputs are being referenced by another CloudFormation stack.

## Output Example:

Create a SSH Security Group as a part of one template  
We create an output that references that security group.

We have the output and we have the name of the Security Group. Then we give a description and

### Outputs:

```
StackSSHSecurityGroup:  
  Description: The SSH Security Group for our Company  
  Value: !Ref MyCompanyWideSSHSecurityGroup  
  Export:  
    Name: SSHSecurityGroup
```

**Value:** we are basically referencing it to the Security Group that is in the Resource and  
**Export:** you can use it to export value name that will be used to reference it in another resource.

## Cross Stack Reference

- For this we use the **Fn::ImportValue** function.
- You can not delete the underlying stack until all the

references are deleted too

### Resources:

```
MySecureInstance:  
  Type: AWS::EC2::Instance  
  Properties:  
    AvailabilityZone: us-east-1a  
    ImageId: ami-a4c7edb2  
    InstanceType: t2.micro  
    SecurityGroups:  
      - !ImportValue SSHSecurityGroup
```

In the security group: you can use **ImportValue** to reuse the Security Group.

# CloudFormation Conditions

- Conditions are used to control the creation of resources or outputs based on a condition.
- Conditions can be whatever you want them to be, but common ones are:
  - Environments
  - AWS Region
  - Any parameter value
- Each condition can reference another condition, parameter value or mapping.

## How to define a condition?

### Conditions:

```
| CreateProdResources: !Equals [ !Ref EnvType, prod ]
```

- The logical ID is for you to choose. It is how you name condition.
- The intrinsic function (logical) can be any of the following:
  - Fn::And
  - Fn::Equals
  - Fn::If
  - Fn::Not
  - Fn::Or

## Using a Condition

- Condition can be applied to resource / outputs / etc..

### Resources:

#### MountPoint:

```
Type: "AWS::EC2::VolumeAttachment"
```

```
Condition: CreateProdResources
```

## CloudFormation Must Know Intrinsic Functions

- Ref
- Fn::GetAtt
- Fn::FindInMap
- Fn::ImportValue
- Fn::Join
- Fn::Sub
- Condition Function (Fn::If, Fn::Not, Fn::Equals, etc..)

### Fn::Ref

- The **Fn::Ref** function can be leverage to reference
- Parameters => returns the value of the parameters
- Resource => returns the physical ID of the underlying resource (Ex:EC2 ID)
- The shorthand for this in YAML is **!Ref**

**DbSubnet1:**

Type: AWS::EC2::Subnet

Properties:

VpcId: !Ref MyVPC

### Fn::GetAtt

- Attributes are attached to any resources you create
- To know the attributes of your resources, the best place to look is the documentation.
- For Example: the AZ of an EC2 machine!

Resources:

  EC2Instance:

    Type: "AWS::EC2::Instance"

    Properties:

      ImageId: ami-1234567

      InstanceType: t2.micro

NewVolume:

  Type: "AWS::EC2::Volume"

  Condition: CreateProdResources

  Properties:

    Size: 100

    AvailabilityZone:

      !GetAtt EC2Instance.AvailabilityZone

### Fn::Sub

Fn::Sub, or !Sub as a shorthand, is used to substitute variables from a text. It is a very handy function that will allow you fully customize your templates.

For Example, you can combine Fn::Sub with Reference or AWS Pseudo variables!

String must contain \${VariableName} and will substitute them

```
!Sub
- String
- { Var1Name: Var1Value, Var2Name: Var2Value }
```

```
!Sub String
```

## User Data in EC2 for CloudFormation

- We can have user data at EC2 instance launch through the console
- We can also include it in CloudFormation
- The important thing to pass is the entire script through the function **Fn::Base64**
- Good to know: user data script log is in **/var/log/cloud-init-output.log**

```
! 3-user-data.yaml x
1 ---
2 Parameters:
3   SSHKey:
4     Type: AWS::EC2::KeyPair::KeyName
5     Description: Name of an existing EC2 KeyPair to enable SSH access to the instance
6
7 Resources:
8   MyInstance:
9     Type: AWS::EC2::Instance
10    Properties:
11      AvailabilityZone: us-east-1a
12      ImageId: ami-009d6802948d06e52
13      InstanceType: t2.micro
14      KeyName: !Ref SSHKey
15      SecurityGroups:
16        - !Ref SSCHttpdSecurityGroup
17      # we install our web server with user data
18      UserData:
19        Fn::Base64: |
20          #!/bin/bash -xe
21          yum update -y
22          yum install -y httpd
23          systemctl start httpd
24          systemctl enable httpd
25          echo "Hello World from user data" > /var/www/html/index.html
26
27      # our EC2 security group
28      SSCHttpdSecurityGroup:
29        Type: AWS::EC2::SecurityGroup
30        Properties:
31          GroupDescription: SSH and HTTP
32          SecurityGroupIngress:
33            - CidrIp: 0.0.0.0/0
34              FromPort: 22
35              IpProtocol: tcp
36              ToPort: 22
37            - CidrIp: 0.0.0.0/0
38              FromPort: 80
39              IpProtocol: tcp
40              ToPort: 80
```

**Fn::Base64** is the script where you have defined the user data and you need to have that "!" pipe in order to make your script work on the line 19, otherwise your script won't work. Our user data installs the simple apache script. And as we did before we need to upload the file and as we did before.

For the SSH key pair you can specify it in the console or if you don't have one, you can create it and put during the creation of the file.

```
~/aws-course ➔ ssh -i CloudFormationKeyPair.pem ec2-user@18.212.193.181
Warning: Permanently added '18.212.193.181' (ECDSA) to the list of known hosts.
```

```
_|_ _|_
_| ( _|_ / Amazon Linux 2 AMI
__\_\_|\_
```

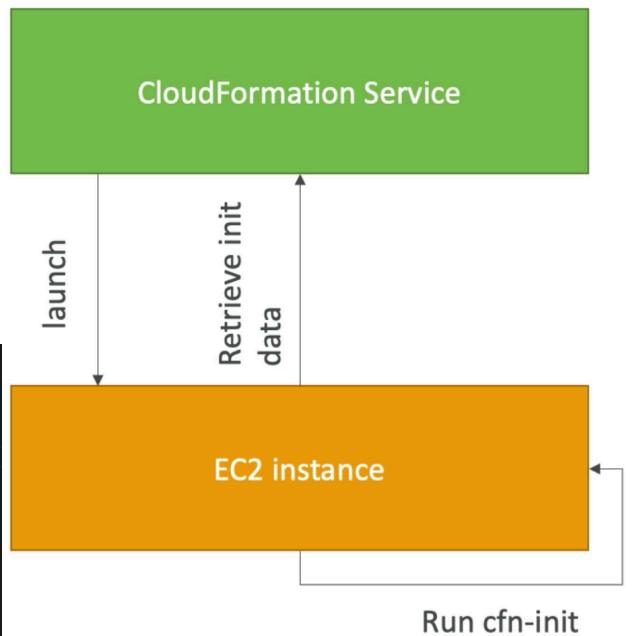
```
https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-172-31-92-82 ~]$ cat /var/log/cloud-init-output.log
```

This is where file for the user data will go good for the debug.

## cfn-init

- AWS::CloudFormation::Init must be in the Metadata of resource
- With the **cfn-init** script, it helps make complex EC2 configuration service to get init data
- Logs go to **/var/log/cfn-init.log**

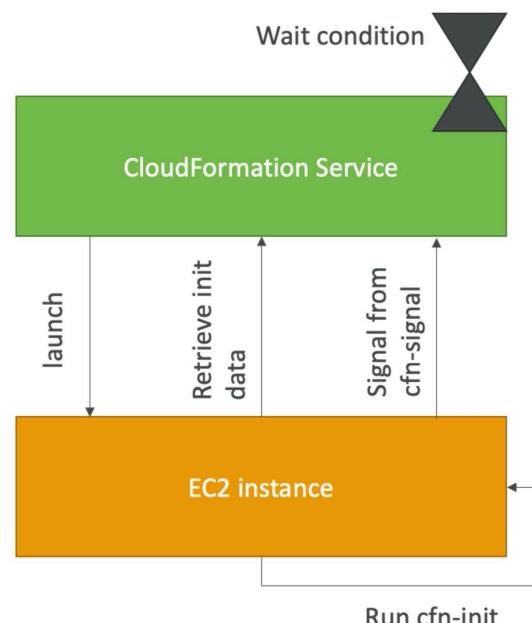
```
! 3-user-data.yaml ! 4-cfn-init.yaml x
1 ---  
2 Parameters:  
3   SSHKey:  
4     Type: AWS::EC2::KeyPair::KeyName  
5     Description: Name of an existing EC2 KeyPair to enable SSH access to the instance  
6  
7 Resources:  
8   MyInstance:  
9     Type: AWS::EC2::Instance  
10    Properties:  
11      AvailabilityZone: us-east-1a  
12      ImageId: ami-009d6802948d06e52  
13      InstanceType: t2.micro  
14      KeyName: !Ref SSHKey  
15      SecurityGroups:  
16        - !Ref SSHSecurityGroup  
17      # we install our web server with user data  
18      UserData:  
19        Fn::Base64:  
20          !Sub |  
21            #!/bin/bash -xe  
22            # Get the latest CloudFormation package  
23            yum update -y aws-cfn-bootstrap  
24            # Start cfn-init  
25            /opt/aws/bin/cfn-init -s ${AWS::StackId} -r MyInstance --region ${AWS::Region} ||  
26            error_exit 'Failed to run cfn-init'  
27  
28      Metadata:  
29        Comment: Install a simple Apache HTTP page  
30        AWS::CloudFormation::Init:  
31          config:  
32            packages:  
33              yum:  
34                httpd: []  
35            files:  
36              "/var/www/html/index.html":  
37                content: |  
38                  <h1>Hello World from EC2 instance!</h1>  
39                  <p>This was created using cfn-init</p>  
40                  File: '000644'  
41            commands:  
42              hello:  
43                command: "echo 'hello world'"  
44            services:  
45              sysvinit:  
46                httpd:  
47                  enabled: 'true'  
48                  ensureRunning: 'true'
```



Ln18: Still a user data  
Ln19: Still a Base64  
Ln20: We are passing !Sub function.  
Ln23: we are updating the **aws-cfn-bootstrap** script.  
Ln25:  
Ln26: Metadata starts  
Ln27: installs apache comment  
Ln28: reps aws cf init block.  
Ln29: config starts, where you start defining entire configuration.  
Ln30: starts with **packages**.  
Install httpd.  
Ln33: **Files** you can define the content inside.  
Ln39: You can define the commands.  
Ln42: You can say what **services** it should start. and the rest is exact same thing.

# Cfn-signal & wait conditions

- We still do not know how to tell CloudFormation that the EC2 instance got properly configured after a **cfn-init**
- For this we can use the **cfn-signal** script!
  - ✓ We run **cfn-signal** right after **cfn-init**
  - ✓ Tell CloudFormation service to keep on going or fail
- We need to define **WaitCondition**:
  - ✓ Block the template until it receives a signal from cfn-signal
  - ✓ We attach a CreationPolicy (also works on EC2, ASG)



```
7 Resources:
8   MyInstance:
9     Type: AWS::EC2::Instance
10    Properties:
11      AvailabilityZone: us-east-1a
12      ImageId: ami-009d6802948d06e52
13      InstanceType: t2.micro
14      KeyName: !Ref SSHKey
15      SecurityGroups:
16        - !Ref SSHSecurityGroup
17      # we install our web server with user data
18      UserData:
19        Fn::Base64:
20          !Sub |
21            #!/bin/bash -xe
22            # Get the latest CloudFormation package
23            yum update -y aws-cfn-bootstrap
24            # Start cfn-init
25            /opt/aws/bin/cfn-init -s ${AWS::StackId} -r MyInstance --region ${AWS::Region}
26            # Start cfn-signal to the wait condition
27            /opt/aws/bin/cfn-signal -e $! --stack ${AWS::StackId} --resource SampleWaitCondition
28            --region ${AWS::Region}
29
30   Metadata:
31     Comment: Install a simple Apache HTTP page
32     AWS::CloudFormation::Init:
33       config:
34         packages:
35           yum:
36             httpd: []
37             files: [
38               "/var/www/html/index.html",
39               content: |
40                 <h1>Hello World from EC2 instance!</h1>
41                 <p>This was created using cfn-init</p>
42                 mode: '000644'
43             commands:
44               hello:
45                 command: "echo 'hello world'"
46             services:
47               sysvinit:
48                 httpd:
49                   enabled: 'true'
49                   ensureRunning: 'true'
50
51   SampleWaitCondition:
52     CreationPolicy:
53       ResourceSignal:
54         Timeout: PT2M
55     Type: AWS::CloudFormation::WaitCondition
56
57     # our EC2 security group
58     SSHSecurityGroup:
59       Type: AWS::EC2::SecurityGroup
59       Properties:
```

We have exit same parameters and exact same user data but on line 27, we have cfn init command to define the conditions and basically it is saying if the line25 will work execute the instance. And on line 27 we are going to signal SampleWaitCondition that is on line50 where we are passing some configuration to the SimpleWaitCondition like how long it should wait.

Line 53: Timeout:PT2M

# CF cfn-signal failures troubleshooting

- Ensure that the AMI you are using has the AWS CF helpers scripts installed. If the AMI doesn't include the helper scripts, you can also download them to your instance.
- Verify that the **cfn-init & cfn-signal** command was successfully run on the instances. You can view logs, such as **/var/log/cloud-init.log** or **/var/log/cfn-init.log** to help you to debug the instance launch.

```
28
29     Metadata:
30         Comment: Install a simple Apache HTTP page
31     AWS::CloudFormation::Init:
32         config:
33             packages:
34                 yum:
35                     httpd: []
36             files:
37                 "/var/www/html/index.html":
38                     content: |
39                         <h1>Hello World from EC2 instance!</h1>
40                         <p>This was created using cfn-init</p>
41                     mode: '000644'
42             commands:
43                 hello:
44                     command: "echo 'boom' && exit 1"
45             services:
46                 sysvinit:
47                     httpd:
48                         enabled: 'true'
49                         ensureRunning: 'true'
```

- You can retrieve the logs by logging in to your instance, but you must disable rollback on failure or else AWS CF deletes the instance after your stack fails to create.

- Verify that the instance has a connection to the internet. If the instance is in VPC the instance should be able to connect to the internet through a NAT device if it is in a private subnet or through an Internet gateway if it is in a public subnet.

Same Script and we will just exit in the command line43 with the exit1 message, that says if the script is not working it will exit with error message.

# CF Rollbacks

## Stack Creation Fails:

- Default: everything rolls back (gets deleted). We can look at the log
- Option to disable rollback and troubleshoot what happened

## Stack Update Fails:

- The stack automatically rolls back to the previous known working state
- Ability to see in the log what happened and error messages.

Timestamp	Logical ID	Status	Status reason
2020-05-06 20:18:57 UTC+0100	FailureDemo	COMPLETE	UPDATE_ROLLBACK_COMPLETE
2020-05-06 20:18:57 UTC+0100	SSHSecurityGroup	COMPLETE	DELETE_COMPLETE
2020-05-06 20:18:56 UTC+0100	ServerSecurityGroup	COMPLETE	DELETE_COMPLETE
2020-05-06 20:18:55 UTC+0100	SSHSecurityGroup	IN_PROGRESS	DELETE_IN_PROGRESS
	ServerSecurity	IN_PROGRESS	DELETE_IN_PROGRESS

But if we disable the rollback it will not go back to the previous condition but it will fail. Sometimes it is good for the troubleshooting. Nothing will be rolled back.

In case if the stack will fail with new update it will roll back to the previous condition. It will not place an update for us. That is what the roll back does for us.

### ▼ Stack creation options

#### Rollback on failure

Specifies whether the stack should be rolled back if stack creation fails.

- Enabled  
 Disabled

#### Timeout

The number of minutes before a stack creation times out.

Minutes

#### Termination protection

Prevents the stack from being accidentally deleted. Once created, you can update this through stack actions.

- Disabled  
 Enabled

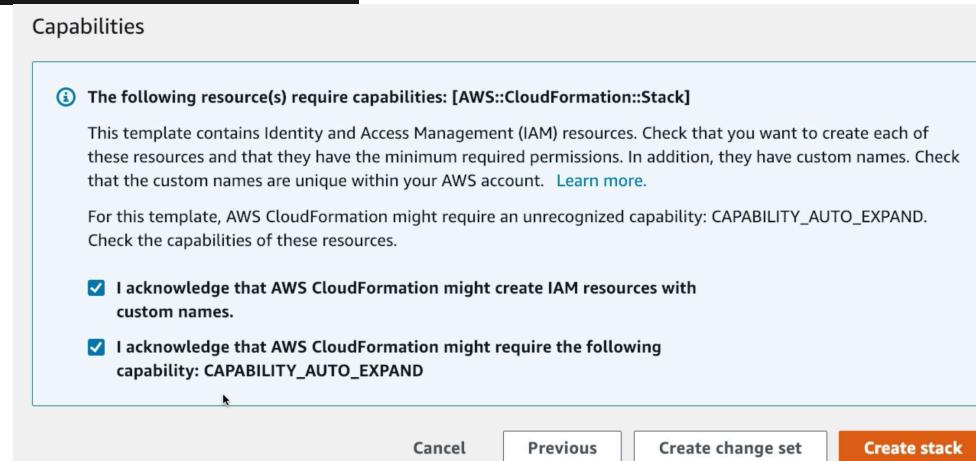
# CF Nested Stacks

- Nested stacks are the stacks as part of other stacks
- They allow you to install isolate repeated patterns / common components in separate stacks and call them from other stacks.
- Example:
  - Load Balancer configuration that is re-used
  - Security Group that is re-used
- Nested Stacks are considered best practice
- To update a nested stack always update the parent (root stack)

```
1 Parameters:
2   SSHKey:
3     Type: AWS::EC2::KeyPair::KeyName
4     Description: Name of an existing EC2 KeyPair to enable SSH access to the instance
5
6 Resources:
7   myStack:
8     Type: AWS::CloudFormation::Stack
9     Properties:
10    TemplateURL:
11      https://s3.amazonaws.com/cloudformation-templates-us-east-1/LAMP_Single_Instance.template
12    Parameters:
13      KeyName: !Ref SSHKey
14      DBName: "mydb"
15      DBUser: "user"
16      DBPassword: "pass"
17      DBRootPassword: "passroot"
18      InstanceType: t2.micro
19      SSHLocation: "0.0.0.0/0"
20
21 Outputs:
22   StackRef:
23     Value: !Ref myStack
24   OutputFromNestedStack:
25     Value: !GetAtt myStack.Outputs.WebsiteURL
```

This is the warning that tells you that you need to have all the proper IAAM rolls to securely deploy the stack.

Nested stacks are basically the file that you can define all the configuration with the resource and upload it with using the stack. For Example: in our case our file that is in S3 bucket has the configuration of files that creates and EC2, DB and other resources.

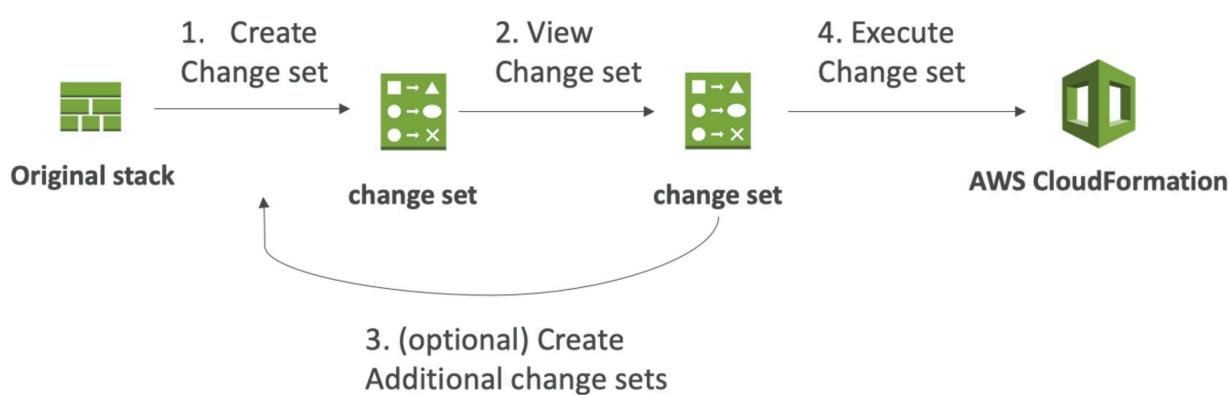


Stacks			
Stack name	Status	Created time	Description
nested-stack-ex-myStack-...	CREATE_COMPLETE	Mon, 17 Dec 2018 11:24:1...	AWS CloudFormation Sam...
nested-stack-ex	CREATE_COMPLETE	Mon, 17 Dec 2018 11:24:0...	-

Warning: never update an child stack, what you usually do is you have to update the parent stack or you can delete the parent stack but never delete the child stack Or you have to delete all the provisioned resources manually

# ChangesSets

- When you update a stack, you need to know what changes before it happens for greater confidence.
- ChangeSets won't say if the update will be successful.



show us the changes that has been made. And this is how our change set looks like. But as you can see as of now we don't have any changeSet lets create it. So there are two ways to create it and one is you can right click on the existing stack or you can press create change stack.

**CloudFormation > Stacks > stak-cs-demo: Change sets**

**stak-cs-demo**

**Change sets**

**Name**      **Created time**      **Status**      **Description**

**Empty change sets**  
No change sets to display

**Create change set**

**CloudFormation > Stacks > stak-cs-demo: Create change set**

**Step 1 Specify template**

**Step 2 Specify stack details**

**Step 3 Configure stack options**

**Step 4 Review**

**Create change set for stak-cs-demo**

**Prerequisite - Prepare template**

**Prepare template**  
Every stack is based on a template. A template is a JSON or YAML file that contains configuration information about the AWS resources you want to include in the stack.

Use current template     Replace current template     Edit template in designer

**CloudFormation > Stacks > stak-cs-demo: Changes**

**Changes (4)**

Action	Logical ID	Physical ID	Resource type	Replacement
Add	MyEIP	-	AWS::EC2::EIP	-
Modify	Myinstance	i-0f7e96d51734fc5be	AWS::EC2::Instance	True
Add	SSHSecurityGroup	-	AWS::EC2::SecurityGroup	-
Add	ServerSecurityGroup	-	AWS::EC2::SecurityGroup	-

We have been using the changeSet so far until now. But to see the changeSet you need to upload the updated template file for that. And it will

We will go and replace the existing stack with the new one. And press create a change set. But our changeSet is not implemented yet. It will create a change but to implement the changes that we had defined in the changeSet we need to press execute.

And we can see the all the changes that will be implemented to our new infrastructure.

# Retain Data on Deletes

You can put DeletePolicy on any resource to control what happens when the CF template is deleted.

## DeletionPolicy=Retain:

- Specify on resource to preserve / backup in case of CF deletes.
- To keep a resource, specify Retain (works for any resource / nested stack)

## DeletionPolicy=Snapshot:

- EBS Volume, ElasticCache Cluster, ElasticCache Replication Group
- RDS DBInstance, RDS DBCluster, Redshift Cluster

## DeletePolicy=Delete (default behavior):

- Note: AWS::RDS::DBCluster resources, the default policy is Snapshot.
- Note: to delete an S3Bucket, you need to first empty the bucket of its content.

```
! 6-cfn-signal-failure.yaml ! 7-nestedstacks.yaml ! 8-changesets.yaml
1 ---
2 Resources:
3   MySG:
4     Type: AWS::EC2::SecurityGroup
5     DeletionPolicy: Retain
6     Properties:
7       GroupDescription: Enable SSH access via port 22
8       SecurityGroupIngress:
9         - CidrIp: 0.0.0.0/0
10      FromPort: 22
11      IpProtocol: tcp
12      ToPort: 22
13
14   MyEBS:
15     Type: AWS::EC2::Volume
16     DeletionPolicy: Snapshot
17     Properties:
18       AvailabilityZone: us-east-1a
19       Size: 1
20       VolumeType: gp2
```

As you can see on the line 5 we have enabled the deletionPolicy to Retain: It should keep our Security Group.

And in the second resource line 16, we have enabled the snapshot of our EBS volume. It should take a snapshot of it.

Timestamp	Logical ID	Status	Status reason
17 Dec 2018 11:47:01	Snapshot-vol-0b77efec64c3978f0	CREATE_COMPLETE	-
17 Dec 2018 11:46:15	MySG	DELETE_SKIPPED	-
17 Dec 2018 11:46:15	Snapshot-vol-0b77efec64c3978f0	CREATE_IN_PROGRESS	-
17 Dec 2018 11:46:14	MyEBS	DELETE_IN_PROGRESS	-
17 Dec 2018 11:46:13	deleteionpolicy	DELETE_IN_PROGRESS	User Initiated

So what we need to remember is we have a snapshot of EBS volume and Security Group existing. Since we specified it in the resource section.

# Termination Protection on Stacks

To prevent accidental deletes of the CF templates use **TerminationProtection**.

The screenshot shows the AWS CloudFormation console with the 'Stacks' tab selected. A stack named 'test' is listed with a status of 'CREATE\_COMPLETE'. A context menu is open over the stack, with the 'Edit termination protection' option highlighted. To the right, a modal dialog titled 'Enable termination protection' asks if the user wants to enable termination protection for the 'test' stack. The current setting is 'Disabled'. A note below states that termination protection prevents the stack from being deleted accidentally. At the bottom of the dialog are 'Cancel' and 'Enable' buttons.

## We will create an AutoScaling Group

```
auto-scaling > cloudformation > ! 0-asg-creation-policy.yml > {} Parameters > {} LatestLinuxAmiId >
1 Parameters:
2   LatestLinuxAmiId:
3     Type: AWS::SSM::Parameter::Value<AWS::EC2::Image::Id>
4     Default: '/aws/service/ami-amazon-linux-latest/amzn2-ami-hvm-x86_64-gp2'
5
6 Resources:
7   AutoScalingGroup:
8     Type: AWS::AutoScaling::AutoScalingGroup
9     Properties:
10    AvailabilityZones:
11      Fn::GetAZs:
12        Ref: "AWS::Region"
13    LaunchConfigurationName:
14      Ref: LaunchConfig
15    DesiredCapacity: '3'
16    MinSize: '1'
17    MaxSize: '4'
18  CreationPolicy:
19    ResourceSignal:
20      Count: '3'
21      Timeout: PT15M
22
23  LaunchConfig:
24    Type: AWS::AutoScaling::LaunchConfiguration
25    Properties:
26      ImageId: !Ref LatestLinuxAmiId
27      InstanceType: t2.micro
28      UserData:
29        Fn::Base64:
30          !Sub |
31            #!/bin/bash -xe
32            yum update -y aws-cfn-bootstrap
33            /opt/aws/bin/cfn-signal -e $? --stack ${AWS::StackName} --resource
AutoScalingGroup --region ${AWS::Region}
```

**CerationPolicy** The creation policy requires three success signals and times out after 15 minutes. To have instances wait for an Elastic Load Balancing health check before they signal success, add a health-check verification by using the cfn-init helper script.